

Introduction

In this lab, we will learn how to use the Keras library to build models for classification problems. We will use the popular MNIST dataset, a dataset of images, for a change.

The MNIST database, short for Modified National Institute of Standards and Technology database, is a large database of handwritten digits that is commonly used for training various image processing systems. The database is also widely used for training and testing in the field of machine learning.

The MNIST database contains 60,000 training images and 10,000 testing images of digits written by high school students and employees of the United States Census Bureau.

Also, this way, will get to compare how conventional neural networks compare to convolutional neural networks, that we will build in the next module.

Objective for this Notebook

1. Use of MNIST database for training various image processing systems 2. Build a Neural Network 3. Train and Test the Network.

This link will be used by your peers to assess your project. In your web app, your peers will be able to upload an image, which will then be classified using your custom classifier you connected to the web app. Your project will be graded by how accurately your app can classify Fire, Smoke and Neutral (No Fire or Smoke).

Table of Contents

Import Keras and Packages

Let's start by importing Keras and some of its modules.

All Libraries required for this lab are listed below. The libraries pre-installed on Skills Network Labs are commented.

If you run this notebook on a different environment, e.g. your desktop, you may need to uncomment and install certain libraries.

```
!pip install numpy==1.21.4
!pip install pandas==1.3.4
!pip install keras==2.1.6
!pip install matplotlib==3.5.0
```

```
import warnings
warnings.filterwarnings('ignore')
```

```
import keras
```

```
from keras.models import Sequential
```

```
from keras.layers import Dense
from keras.utils import to_categorical
```

Using TensorFlow backend.

Since we are dealing with images, let's also import the Matplotlib scripting layer in order to view the images.

```
import matplotlib.pyplot as plt
```

The Keras library conveniently includes the MNIST dataset as part of its API. You can check other datasets within the Keras library [here](#).

So, let's load the MNIST dataset from the Keras library. The dataset is readily divided into a training set and a test set.

```
# import the data
from keras.datasets import mnist
```

```
# read the data
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

```
Downloading data from https://s3.amazonaws.com/img-datasets/mnist.npz
11493376/11490434 [=====] - 0s 0us/step
```

Let's confirm the number of images in each set. According to the dataset's documentation, we should have 60000 images in X_train and 10000 images in the X_test.

```
X_train.shape
(60000, 28, 28)
```

The first number in the output tuple is the number of images, and the other two numbers are the size of the images in dataset. So, each image is 28 pixels by 28 pixels.

Let's visualize the first image in the training set using Matplotlib's scripting layer.

```
plt.imshow(X_train[0])
```

With conventional neural networks, we cannot feed in the image as input as is. So we need to flatten the images into one-dimensional vectors, each of size $1 \times (28 \times 28) = 1 \times 784$.

```
# flatten images into one-dimensional vector
```

```
num_pixels = X_train.shape[1] * X_train.shape[2] # find size of one-dimensional vector
```

```
X_train = X_train.reshape(X_train.shape[0],
num_pixels).astype('float32') # flatten training images
X_test = X_test.reshape(X_test.shape[0], num_pixels).astype('float32')
# flatten test images
```

Since pixel values can range from 0 to 255, let's normalize the vectors to be between 0 and 1.

```
# normalize inputs from 0-255 to 0-1
X_train = X_train / 255
X_test = X_test / 255
```

Finally, before we start building our model, remember that for classification we need to divide our target variable into categories. We use the `to_categorical` function from the Keras Utilities package.

```
# one hot encode outputs
y_train = to_categorical(y_train)
y_test = to_categorical(y_test)
```

```
num_classes = y_test.shape[1]
print(num_classes)
```

```
10
```

Build a Neural Network

```
# define classification model
def classification_model():
    # create model
    model = Sequential()
    model.add(Dense(num_pixels, activation='relu',
input_shape=(num_pixels,)))
    model.add(Dense(100, activation='relu'))
    model.add(Dense(num_classes, activation='softmax'))

    # compile model
    model.compile(optimizer='adam', loss='categorical_crossentropy',
metrics=['accuracy'])
    return model
```

Train and Test the Network

```
# build the model
model = classification_model()

# fit the model
model.fit(X_train, y_train, validation_data=(X_test, y_test),
epochs=10, verbose=2)

# evaluate the model
scores = model.evaluate(X_test, y_test, verbose=0)
```

WARNING:tensorflow:From
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages/keras/
backend/tensorflow_backend.py:68: The name tf.get_default_graph is
deprecated. Please use tf.compat.v1.get_default_graph instead.

WARNING:tensorflow:From
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages/keras/
backend/tensorflow_backend.py:508: The name tf.placeholder is
deprecated. Please use tf.compat.v1.placeholder instead.

WARNING:tensorflow:From
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages/keras/
backend/tensorflow_backend.py:3837: The name tf.random_uniform is
deprecated. Please use tf.random.uniform instead.

WARNING:tensorflow:From
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages/keras/
optimizers.py:757: The name tf.train.Optimizer is deprecated. Please
use tf.compat.v1.train.Optimizer instead.

WARNING:tensorflow:From
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages/keras/
backend/tensorflow_backend.py:3014: The name tf.log is deprecated.
Please use tf.math.log instead.

WARNING:tensorflow:From
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages/tensorf
low/python/ops/math_grad.py:1250:
add_dispatch_support.<locals>.wrapper (from
tensorflow.python.ops.array_ops) is deprecated and will be removed in
a future version.
Instructions for updating:
Use tf.where in 2.0, which has the same broadcast rule as np.where
WARNING:tensorflow:From
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages/keras/
backend/tensorflow_backend.py:977: The name tf.assign_add is
deprecated. Please use tf.compat.v1.assign_add instead.

Train on 60000 samples, validate on 10000 samples
Epoch 1/10

2022-10-19 14:47:21.720078: I
tensorflow/core/platform/cpu_feature_guard.cc:142] Your CPU supports
instructions that this TensorFlow binary was not compiled to use:
SSE4.1 SSE4.2 AVX AVX2 FMA
2022-10-19 14:47:21.736017: I
tensorflow/core/platform/profile_utils/cpu_utils.cc:94] CPU Frequency:
2095145000 Hz
2022-10-19 14:47:21.737425: I
tensorflow/compiler/xla/service/service.cc:168] XLA service

```
0x563419d11080 executing computations on platform Host. Devices:
2022-10-19 14:47:21.737507: I
tensorflow/compiler/xla/service/service.cc:175]   StreamExecutor
device (0): <undefined>, <undefined>
2022-10-19 14:47:21.916240: W
tensorflow/compiler/jit/mark_for_compilation_pass.cc:1412] (One-time
warning): Not using XLA:CPU for cluster because envvar TF_XLA_FLAGS=--
tf_xla_cpu_global_jit was not set. If you want XLA:CPU, either set
that envvar, or use experimental_jit_scope to enable XLA:CPU. To
confirm that XLA is active, pass --vmodule=xla_compilation_cache=1 (as
a proper command-line flag, not via TF_XLA_FLAGS) or set the envvar
XLA_FLAGS=--xla_hlo_profile.
```

```
- 60s - loss: 0.1829 - acc: 0.9442 - val_loss: 0.0918 - val_acc:
0.9728
```

```
Epoch 2/10
```

```
- 61s - loss: 0.0782 - acc: 0.9756 - val_loss: 0.0860 - val_acc:
0.9733
```

```
Epoch 3/10
```

```
- 60s - loss: 0.0543 - acc: 0.9825 - val_loss: 0.0821 - val_acc:
0.9753
```

```
Epoch 4/10
```

```
- 60s - loss: 0.0386 - acc: 0.9871 - val_loss: 0.0727 - val_acc:
0.9772
```

```
Epoch 5/10
```

```
- 60s - loss: 0.0324 - acc: 0.9896 - val_loss: 0.0810 - val_acc:
0.9767
```

```
Epoch 6/10
```

```
- 59s - loss: 0.0285 - acc: 0.9903 - val_loss: 0.0727 - val_acc:
0.9808
```

```
Epoch 7/10
```

```
- 59s - loss: 0.0221 - acc: 0.9927 - val_loss: 0.0840 - val_acc:
0.9804
```

```
Epoch 8/10
```

```
- 59s - loss: 0.0204 - acc: 0.9933 - val_loss: 0.0868 - val_acc:
0.9785
```

```
Epoch 9/10
```

```
- 60s - loss: 0.0171 - acc: 0.9948 - val_loss: 0.0761 - val_acc:
0.9812
```

```
Epoch 10/10
```

```
- 61s - loss: 0.0172 - acc: 0.9945 - val_loss: 0.0931 - val_acc:
0.9798
```

Let's print the accuracy and the corresponding error.

```
print('Accuracy: {}% \n Error: {}'.format(scores[1], 1 - scores[1]))
```

```
Accuracy: 0.9798%
```

```
Error: 0.020199999999999996
```

Just running 10 epochs could actually take over 20 minutes. But enjoy the results as they are getting generated.

Sometimes, you cannot afford to retrain your model everytime you want to use it, especially if you are limited on computational resources and training your model can take a long time. Therefore, with the Keras library, you can save your model after training. To do that, we use the save method.

```
model.save('classification_model.h5')
```

Since our model contains multidimensional arrays of data, then models are usually saved as .h5 files.

When you are ready to use your model again, you use the load_model function from keras.models.

```
from keras.models import load_model  
pretrained_model = load_model('classification_model.h5')
```

Thank you for completing this lab!

This notebook was created by [Alex Aklson](#). I hope you found this lab interesting and educational. Feel free to contact me if you have any questions!

Change Log

Date (YYYY-MM-DD)	Version	Changed By	Change Description
2020-09-21	2.0	Srishti	Migrated Lab to Markdown and added to course repo in GitLab

© IBM Corporation 2020. All rights reserved.

This notebook is part of a course on **Coursera** called *Introduction to Deep Learning & Neural Networks with Keras*. If you accessed this notebook outside the course, you can take this course online by clicking [here](#).

Copyright © 2019 [IBM Developer Skills Network](#). This notebook and its source code are released under the terms of the [MIT License](#).