

In this lab, we will learn how to use the Keras library to build convolutional neural networks. We will also use the popular MNIST dataset and we will compare our results to using a conventional neural network.

Objective for this Notebook

1. How to use the Keras library to build convolutional neural networks. 2. Convolutional Neural Network with One Convolutional and Pooling Layers. 3. Convolutional Neural Network with Two Convolutional and Pooling Layers.

## Table of Contents

### Import Keras and Packages

Let's start by importing the keras libraries and the packages that we would need to build a neural network.

*# All Libraries required for this lab are listed below. The libraries pre-installed on Skills Network Labs are commented.  
# If you run this notebook on a different environment, e.g. your desktop, you may need to uncomment and install certain libraries.*

```
!pip install numpy==1.21.4  
!pip install pandas==1.3.4  
!pip install keras==2.1.6
```

Collecting numpy==1.21.4

Downloading numpy-1.21.4-cp37-cp37m-manylinux\_2\_12\_x86\_64.manylinux2010\_x86\_64.whl (15.7 MB)  
15.7/15.7 MB 59.9 MB/s eta

0:00:0000:0100:01

py

Attempting uninstall: numpy

Found existing installation: numpy 1.21.6

Uninstalling numpy-1.21.6:

Successfully uninstalled numpy-1.21.6

Successfully installed numpy-1.21.4

Collecting pandas==1.3.4

Downloading pandas-1.3.4-cp37-cp37m-manylinux\_2\_17\_x86\_64.manylinux2014\_x86\_64.whl (11.3 MB)  
11.3/11.3 MB 64.9 MB/s eta

0:00:0000:0100:01

Requirement already satisfied: python-dateutil>=2.7.3 in

/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from pandas==1.3.4) (2.8.2)

Requirement already satisfied: pytz>=2017.3 in

/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from

```
pandas==1.3.4) (2022.2.1)
Requirement already satisfied: numpy>=1.17.3 in
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from
pandas==1.3.4) (1.21.4)
Requirement already satisfied: six>=1.5 in
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from
python-dateutil>=2.7.3->pandas==1.3.4) (1.16.0)
Installing collected packages: pandas
  Attempting uninstall: pandas
    Found existing installation: pandas 1.3.5
    Uninstalling pandas-1.3.5:
      Successfully uninstalled pandas-1.3.5
Successfully installed pandas-1.3.4
Requirement already satisfied: keras==2.1.6 in
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (2.1.6)
Requirement already satisfied: six>=1.9.0 in
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from
keras==2.1.6) (1.16.0)
Requirement already satisfied: pyyaml in
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from
keras==2.1.6) (6.0)
Requirement already satisfied: h5py in
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from
keras==2.1.6) (2.8.0)
Requirement already satisfied: numpy>=1.9.1 in
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from
keras==2.1.6) (1.21.4)
Requirement already satisfied: scipy>=0.14 in
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from
keras==2.1.6) (1.7.3)
```

```
import warnings
warnings.filterwarnings('ignore')
```

```
import keras
from keras.models import Sequential
from keras.layers import Dense
from keras.utils import to_categorical
```

When working with convolutional neural networks in particular, we will need additional packages.

```
from keras.layers.convolutional import Conv2D # to add convolutional
layers
from keras.layers.convolutional import MaxPooling2D # to add pooling
layers
from keras.layers import Flatten # to flatten data for fully connected
layers
```

## Convolutional Layer with One set of convolutional and pooling layers

```
# import data
from keras.datasets import mnist

# load data
(X_train, y_train), (X_test, y_test) = mnist.load_data()

# reshape to be [samples][pixels][width][height]
X_train = X_train.reshape(X_train.shape[0], 28, 28,
1).astype('float32')
X_test = X_test.reshape(X_test.shape[0], 28, 28, 1).astype('float32')
```

Downloading data from <https://s3.amazonaws.com/img-datasets/mnist.npz>  
11493376/11490434 [=====] - 0s 0us/step

Let's normalize the pixel values to be between 0 and 1

```
X_train = X_train / 255 # normalize training data
X_test = X_test / 255 # normalize test data
```

Next, let's convert the target variable into binary categories

```
y_train = to_categorical(y_train)
y_test = to_categorical(y_test)
```

```
num_classes = y_test.shape[1] # number of categories
```

Next, let's define a function that creates our model. Let's start with one set of convolutional and pooling layers.

```
def convolutional_model():

    # create model
    model = Sequential()
    model.add(Conv2D(16, (5, 5), strides=(1, 1), activation='relu',
input_shape=(28, 28, 1)))
    model.add(MaxPooling2D(pool_size=(2, 2), strides=(2, 2)))

    model.add(Flatten())
    model.add(Dense(100, activation='relu'))
    model.add(Dense(num_classes, activation='softmax'))

    # compile model
    model.compile(optimizer='adam', loss='categorical_crossentropy',
metrics=['accuracy'])
    return model
```

Finally, let's call the function to create the model, and then let's train it and evaluate it.

```
# build the model
model = convolutional_model()
```

```
# fit the model
model.fit(X_train, y_train, validation_data=(X_test, y_test),
epochs=10, batch_size=200, verbose=2)
```

```
# evaluate the model
scores = model.evaluate(X_test, y_test, verbose=0)
print("Accuracy: {} \n Error: {}".format(scores[1], 100-
scores[1]*100))
```

```
WARNING:tensorflow:From
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages/keras/
backend/tensorflow_backend.py:68: The name tf.get_default_graph is
deprecated. Please use tf.compat.v1.get_default_graph instead.
```

```
WARNING:tensorflow:From
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages/keras/
backend/tensorflow_backend.py:508: The name tf.placeholder is
deprecated. Please use tf.compat.v1.placeholder instead.
```

```
WARNING:tensorflow:From
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages/keras/
backend/tensorflow_backend.py:3837: The name tf.random_uniform is
deprecated. Please use tf.random.uniform instead.
```

```
WARNING:tensorflow:From
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages/keras/
backend/tensorflow_backend.py:3661: The name tf.nn.max_pool is
deprecated. Please use tf.nn.max_pool2d instead.
```

```
WARNING:tensorflow:From
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages/keras/
optimizers.py:757: The name tf.train.Optimizer is deprecated. Please
use tf.compat.v1.train.Optimizer instead.
```

```
WARNING:tensorflow:From
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages/keras/
backend/tensorflow_backend.py:3014: The name tf.log is deprecated.
Please use tf.math.log instead.
```

```
WARNING:tensorflow:From
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages/tensorf
low/python/ops/math_grad.py:1250:
add_dispatch_support.<locals>.wrapper (from
tensorflow.python.ops.array_ops) is deprecated and will be removed in
a future version.
```

```
Instructions for updating:
Use tf.where in 2.0, which has the same broadcast rule as np.where
```

```
WARNING:tensorflow:From
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages/keras/
```

backend/tensorflow\_backend.py:977: The name tf.assign\_add is deprecated. Please use tf.compat.v1.assign\_add instead.

Train on 60000 samples, validate on 10000 samples  
Epoch 1/10

2022-10-19 16:51:53.074642: I  
tensorflow/core/platform/cpu\_feature\_guard.cc:142] Your CPU supports instructions that this TensorFlow binary was not compiled to use: SSE4.1 SSE4.2 AVX AVX2 AVX512F FMA  
2022-10-19 16:51:53.084213: I  
tensorflow/core/platform/profile\_utils/cpu\_utils.cc:94] CPU Frequency: 2593905000 Hz  
2022-10-19 16:51:53.085137: I  
tensorflow/compiler/xla/service/service.cc:168] XLA service 0x562d0c37d330 executing computations on platform Host. Devices:  
2022-10-19 16:51:53.085210: I  
tensorflow/compiler/xla/service/service.cc:175] StreamExecutor device (0): <undefined>, <undefined>  
2022-10-19 16:51:53.195225: W  
tensorflow/compiler/jit/mark\_for\_compilation\_pass.cc:1412] (One-time warning): Not using XLA:CPU for cluster because envvar TF\_XLA\_FLAGS=--tf\_xla\_cpu\_global\_jit was not set. If you want XLA:CPU, either set that envvar, or use experimental\_jit\_scope to enable XLA:CPU. To confirm that XLA is active, pass --vmodule=xla\_compilation\_cache=1 (as a proper command-line flag, not via TF\_XLA\_FLAGS) or set the envvar XLA\_FLAGS=--xla\_hlo\_profile.

- 41s - loss: 0.2678 - acc: 0.9293 - val\_loss: 0.0892 - val\_acc: 0.9716

Epoch 2/10

- 39s - loss: 0.0765 - acc: 0.9782 - val\_loss: 0.0546 - val\_acc: 0.9831

Epoch 3/10

- 39s - loss: 0.0537 - acc: 0.9843 - val\_loss: 0.0463 - val\_acc: 0.9857

Epoch 4/10

- 38s - loss: 0.0422 - acc: 0.9876 - val\_loss: 0.0435 - val\_acc: 0.9852

Epoch 5/10

- 39s - loss: 0.0337 - acc: 0.9903 - val\_loss: 0.0448 - val\_acc: 0.9854

Epoch 6/10

- 38s - loss: 0.0278 - acc: 0.9923 - val\_loss: 0.0394 - val\_acc: 0.9863

Epoch 7/10

- 39s - loss: 0.0228 - acc: 0.9933 - val\_loss: 0.0434 - val\_acc: 0.9865

Epoch 8/10

- 38s - loss: 0.0194 - acc: 0.9940 - val\_loss: 0.0378 - val\_acc:

```
0.9885
Epoch 9/10
- 38s - loss: 0.0162 - acc: 0.9953 - val_loss: 0.0458 - val_acc:
0.9872
Epoch 10/10
- 39s - loss: 0.0122 - acc: 0.9965 - val_loss: 0.0360 - val_acc:
0.9878
Accuracy: 0.9878
Error: 1.2199999999999999
```

---

## Convolutional Layer with two sets of convolutional and pooling layers

Let's redefine our convolutional model so that it has two convolutional and pooling layers instead of just one layer of each.

```
def convolutional_model():

    # create model
    model = Sequential()
    model.add(Conv2D(16, (5, 5), activation='relu', input_shape=(28,
28, 1)))
    model.add(MaxPooling2D(pool_size=(2, 2), strides=(2, 2)))

    model.add(Conv2D(8, (2, 2), activation='relu'))
    model.add(MaxPooling2D(pool_size=(2, 2), strides=(2, 2)))

    model.add(Flatten())
    model.add(Dense(100, activation='relu'))
    model.add(Dense(num_classes, activation='softmax'))

    # Compile model
    model.compile(optimizer='adam', loss='categorical_crossentropy',
metrics=['accuracy'])
    return model
```

Now, let's call the function to create our new convolutional neural network, and then let's train it and evaluate it.

```
# build the model
model = convolutional_model()

# fit the model
model.fit(X_train, y_train, validation_data=(X_test, y_test),
epochs=10, batch_size=200, verbose=2)

# evaluate the model
scores = model.evaluate(X_test, y_test, verbose=0)
```

```
print("Accuracy: {} \n Error: {}".format(scores[1], 100-  
scores[1]*100))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/10

- 41s - loss: 0.4767 - acc: 0.8632 - val\_loss: 0.1346 - val\_acc:  
0.9630

Epoch 2/10

- 41s - loss: 0.1216 - acc: 0.9636 - val\_loss: 0.0820 - val\_acc:  
0.9756

Epoch 3/10

- 41s - loss: 0.0841 - acc: 0.9742 - val\_loss: 0.0612 - val\_acc:  
0.9818

Epoch 4/10

- 42s - loss: 0.0675 - acc: 0.9793 - val\_loss: 0.0568 - val\_acc:  
0.9812

Epoch 5/10

- 41s - loss: 0.0571 - acc: 0.9827 - val\_loss: 0.0526 - val\_acc:  
0.9834

Epoch 6/10

- 41s - loss: 0.0515 - acc: 0.9838 - val\_loss: 0.0434 - val\_acc:  
0.9859

Epoch 7/10

- 41s - loss: 0.0448 - acc: 0.9862 - val\_loss: 0.0370 - val\_acc:  
0.9893

Epoch 8/10

- 41s - loss: 0.0397 - acc: 0.9875 - val\_loss: 0.0382 - val\_acc:  
0.9881

Epoch 9/10

- 41s - loss: 0.0372 - acc: 0.9885 - val\_loss: 0.0363 - val\_acc:  
0.9877

Epoch 10/10

**Thank you for completing this lab!**

This notebook was created by [Alex Aklson](#). I hope you found this lab interesting and educational. Feel free to contact me if you have any questions!

## Change Log

Date (YYYY-MM-DD)	Version	Changed By	Change Description
2020-09-21	2.0	Srishti	Migrated Lab to Markdown and added to course repo in GitLab

**© IBM Corporation 2020. All rights reserved.**

This notebook is part of a course on **Coursera** called *Introduction to Deep Learning & Neural Networks with Keras*. If you accessed this notebook outside the course, you can take this course online by clicking [here](#).

Copyright © 2019 [IBM Developer Skills Network](#). This notebook and its source code are released under the terms of the [MIT License](#).