

Hypothesis_Testing

May 19, 2022

1 Hypothesis Testing

Estimated time needed: **30** minutes

The goal of hypothesis testing is to answer the question, “Given a sample and an apparent effect, what is the probability of seeing such an effect by chance?” The first step is to quantify the size of the apparent effect by choosing a test statistic (t-test, ANOVA, etc). The next step is to define a null hypothesis, which is a model of the system based on the assumption that the apparent effect is not real. Then compute the p-value, which is the probability of the null hypothesis being true, and finally interpret the result of the p-value, if the value is low, the effect is said to be statistically significant, which means that the null hypothesis may not be accurate.

1.1 Objectives

- Import Libraries
 - Lab exercises
 - Stating the hypothesis
 - Levene’s Test for equality
 - Preparing your data for hypothesis testing
 - Quiz
-

1.2 Import Libraries

All Libraries required for this lab are listed below. The libraries pre-installed on Skills Network Labs are commented. If you run this notebook in a different environment, e.g. your desktop, you may need to uncomment and install certain libraries.

```
[ ]: #install specific version of libraries used in lab  
#! mamba install pandas==1.3.3  
#! mamba install numpy=1.21.2  
#! mamba install scipy=1.7.1-y  
#! mamba install seaborn=0.9.0-y  
#! mamba install matplotlib=3.4.3-y  
#! mamba install statsmodels=0.12.0-y
```

Import the libraries we need for the lab

```
[1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import scipy.stats
```

Read in the csv file from the URL using the request library

```
[2]: ratings_url = 'https://cf-courses-data.s3.us.cloud-object-storage.appdomain.
↳cloud/IBMDDeveloperSkillsNetwork-ST0151EN-SkillsNetwork/labs/teachingratings.
↳csv'
ratings_df = pd.read_csv(ratings_url)
```

```
[3]: ratings_df.head()
```

```
[3]: minority  age  gender  credits    beauty  eval  division  native  tenure  \
0         yes   36  female    more  0.289916   4.3    upper    yes    yes
1         yes   36  female    more  0.289916   3.7    upper    yes    yes
2         yes   36  female    more  0.289916   3.6    upper    yes    yes
3         yes   36  female    more  0.289916   4.4    upper    yes    yes
4         no    59   male    more -0.737732   4.5    upper    yes    yes

      students  allstudents  prof  PrimaryLast  vismin  female  single_credit  \
0          24          43     1           0        1        1              0
1          86         125     1           0        1        1              0
2          76         125     1           0        1        1              0
3          77         123     1           1        1        1              0
4          17          20     2           0        0        0              0

      upper_division  English_speaker  tenured_prof
0                 1                 1             1
1                 1                 1             1
2                 1                 1             1
3                 1                 1             1
4                 1                 1             1
```

1.3 Lab Exercises

1.3.1 T-Test: Using the teachers' rating data set, does gender affect teaching evaluation rates?

We will be using the t-test for independent samples. For the independent t-test, the following assumptions must be met.

- One independent, categorical variable with two levels or group
- One dependent continuous variable
- Independence of the observations. Each subject should belong to only one group. There is no relationship between the observations in each group.
- The dependent variable must follow a normal distribution

- Assumption of homogeneity of variance


State the hypothesis

- $H_0: \mu_1 = \mu_2$ (“there is no difference in evaluation scores between male and females”)
- $H_1: \mu_1 \neq \mu_2$ (“there is a difference in evaluation scores between male and females”)

We can plot the dependent variable with a histogram

```
[3]: ax = sns.distplot(ratings_df['eval'],
                        bins=20,
                        kde=True,
                        color='red',
                        hist_kws={"linewidth": 15, 'alpha': 1})
ax.set(xlabel='Normal Distribution', ylabel='Frequency')
## we can assume it is normal
```

```
[3]: [Text(0.5, 0, 'Normal Distribution'), Text(0, 0.5, 'Frequency')]
```



We can use the Levene’s Test in Python to check test significance

```
[4]: scipy.stats.levene(ratings_df[ratings_df['gender'] == 'female']['eval'],
                        ratings_df[ratings_df['gender'] == 'male']['eval'],
                        center='mean')
# since the p-value is greater than 0.05 we can assume equality of variance
```

```
[4]: LeveneResult(statistic=0.19032922435292574, pvalue=0.6628469836244741)
```

Use the ttest_ind from the scipy_stats library

```
[5]: scipy.stats.ttest_ind(ratings_df[ratings_df['gender'] == 'female']['eval'],
                           ratings_df[ratings_df['gender'] == 'male']['eval'],
                           equal_var = True)
```

```
[5]: Ttest_indResult(statistic=-3.249937943510772, pvalue=0.0012387609449522217)
```

Conclusion: Since the p-value is less than alpha value 0.05, we reject the null hypothesis as there is enough proof that there is a statistical difference in teaching evaluations based on gender

1.3.2 ANOVA: Using the teachers' rating data set, does beauty score for instructors differ by age?

First, we group the data into categories as the one-way ANOVA can't work with continuous variable - using the example from the video, we will create a new column for this newly assigned group our categories will be teachers that are:

- 40 years and younger
- between 40 and 57 years
- 57 years and older

```
[4]: ratings_df.loc[(ratings_df['age'] <= 40), 'age_group'] = '40 years and younger'
ratings_df.loc[(ratings_df['age'] > 40)&(ratings_df['age'] < 57), 'age_group'] =
    'between 40 and 57 years'
ratings_df.loc[(ratings_df['age'] >= 57), 'age_group'] = '57 years and older'
ratings_df.head()
```

```
[4]: minority age gender credits beauty eval division native tenure \
0      yes  36 female    more  0.289916  4.3    upper    yes    yes
1      yes  36 female    more  0.289916  3.7    upper    yes    yes
2      yes  36 female    more  0.289916  3.6    upper    yes    yes
3      yes  36 female    more  0.289916  4.4    upper    yes    yes
4       no  59   male    more -0.737732  4.5    upper    yes    yes

students allstudents prof PrimaryLast vismin female single_credit \
0         24          43    1           0        1         1         0
1         86         125    1           0        1         1         0
2         76         125    1           0        1         1         0
3         77         123    1           1        1         1         0
4         17          20    2           0        0         0         0

upper_division English_speaker tenured_prof age_group
0              1              1            1 40 years and younger
1              1              1            1 40 years and younger
2              1              1            1 40 years and younger
3              1              1            1 40 years and younger
```

4 1 1 1 57 years and older

State the hypothesis

- $H_0 : \mu_1 = \mu_2 = \mu_3$ (the three population means are equal)
- H_1 : At least one of the means differ

Test for equality of variance

```
[7]: scipy.stats.levene(ratings_df[ratings_df['age_group'] == '40 years and
↳younger']['beauty'],
                        ratings_df[ratings_df['age_group'] == 'between 40 and 57
↳years']['beauty'],
                        ratings_df[ratings_df['age_group'] == '57 years and
↳older']['beauty'],
                        center='mean')
# since the p-value is less than 0.05, the variance are not equal, for the
↳purposes of this exercise, we will move along
```

```
[7]: LeveneResult(statistic=8.60005668392584, pvalue=0.000215366180993476)
```

First, separate the three samples (one for each job category) into a variable each.

```
[8]: forty_lower = ratings_df[ratings_df['age_group'] == '40 years and
↳younger']['beauty']
forty_fiftyseven = ratings_df[ratings_df['age_group'] == 'between 40 and 57
↳years']['beauty']
fiftyseven_older = ratings_df[ratings_df['age_group'] == '57 years and
↳older']['beauty']
```

Now, run a one-way ANOVA.

```
[9]: f_statistic, p_value = scipy.stats.f_oneway(forty_lower, forty_fiftyseven,
↳fiftyseven_older)
print("F_Statistic: {0}, P-Value: {1}".format(f_statistic,p_value))
```

```
F_Statistic: 17.597558611010122, P-Value: 4.3225489816137975e-08
```

Conclusion: Since the p-value is less than 0.05, we will reject the null hypothesis as there is significant evidence that at least one of the means differ.

1.3.3 ANOVA: Using the teachers' rating data set, does teaching evaluation score for instructors differ by age?

Test for equality of variance

```
[10]: scipy.stats.levene(ratings_df[ratings_df['age_group'] == '40 years and
↳younger']['eval'],
                        ratings_df[ratings_df['age_group'] == 'between 40 and 57
↳years']['eval'],
```

```
ratings_df[ratings_df['age_group'] == '57 years and_
↳older']['eval'],
center='mean')
```

```
[10]: LeveneResult(statistic=3.820237661494229, pvalue=0.02262141852021939)
```

```
[11]: forty_lower_eval = ratings_df[ratings_df['age_group'] == '40 years and_
↳younger']['eval']
forty_fiftyseven_eval = ratings_df[ratings_df['age_group'] == 'between 40 and_
↳57 years']['eval']
fiftyseven_older_eval = ratings_df[ratings_df['age_group'] == '57 years and_
↳older']['eval']
```

```
[12]: f_statistic, p_value = scipy.stats.f_oneway(forty_lower_eval,
↳forty_fiftyseven_eval, fiftyseven_older_eval)
print("F_Statistic: {0}, P-Value: {1}".format(f_statistic,p_value))
```

```
F_Statistic: 1.2226327996572206, P-Value: 0.29540894225417536
```

Conclusion: Since the p-value is greater than 0.05, we will fail to reject the null hypothesis as there is no significant evidence that at least one of the means differ.

1.3.4 Chi-square: Using the teachers' rating data set, is there an association between tenure and gender?

State the hypothesis:

- H_0 : The proportion of teachers who are tenured is independent of gender
- H_1 : The proportion of teachers who are tenured is associated with gender

Create a Cross-tab table

```
[13]: cont_table = pd.crosstab(ratings_df['tenure'], ratings_df['gender'])
cont_table
```

```
[13]: gender  female  male
tenure
no           50     52
yes          145    216
```

Use the scipy.stats library and set correction equals False as that will be the same answer when done by hand, it returns: 2 value, p-value, degree of freedom, and expected values.

```
[14]: scipy.stats.chi2_contingency(cont_table, correction = True)
```

```
[14]: (2.20678166999886,
0.1374050603563787,
1,
array([[ 42.95896328,  59.04103672],
```

```
[152.04103672, 208.95896328]]))
```

Conclusion: Since the p-value is greater than 0.05, we fail to reject the null hypothesis. As there is no sufficient evidence that teachers are tenured as a result of gender.


1.3.5 Correlation: Using the teachers rating dataset, Is teaching evaluation score correlated with beauty score?

State the hypothesis:

- H_0 : Teaching evaluation score is not correlated with beauty score
- H_1 : Teaching evaluation score is correlated with beauty score

Since they are both continuous variables we can use a pearson correlation test and draw a scatter plot

```
[15]: ax = sns.scatterplot(x="beauty", y="eval", data=ratings_df)
```



```
[16]: scipy.stats.pearsonr(ratings_df['beauty'], ratings_df['eval'])
```

```
[16]: (0.1890390908404521, 4.247115419812614e-05)
```

Conclusion: Since the p-value (Sig. (2-tailed) < 0.05 , we reject the Null hypothesis and conclude that there exists a relationship between beauty and teaching evaluation score.

1.4 Practice Questions

1.4.1 Question 1: Using the teachers rating data set, does tenure affect teaching evaluation scores?

- Use $\alpha = 0.05$

```
[17]: ## insert code here
      scipy.stats.ttest_ind(ratings_df[ratings_df['tenure'] == 'yes']['eval'],
                           ratings_df[ratings_df['tenure'] == 'no']['eval'], equal_var_
                           ↪= True)
```

```
[17]: Ttest_indResult(statistic=-2.8046798258451777, pvalue=0.005249471210198792)
```

Double-click [here](#) for the solution.

1.4.2 Question 2: Using the teachers rating data set, is there an association between age and tenure?

- Discretize the age into three groups 40 years and younger, between 40 and 57 years, 57 years and older (This has already been done for you above.)
- What is your conclusion at $\alpha = 0.01$ and $\alpha = 0.05$?

```
[18]: ## insert code here
      scipy.stats.chi2_contingency(cont_table, correction = True)
```

```
[18]: (2.20678166999886,
      0.1374050603563787,
      1,
      array([[ 42.95896328,  59.04103672],
             [152.04103672, 208.95896328]]))
```

Double-click [here](#) for a hint.

Double-click [here](#) for the solution.

1.4.3 Question 3: Test for equality of variance for beauty scores between tenured and non-tenured instructors

- Use $\alpha = 0.05$

```
[19]: ## insert code here
      scipy.stats.levene(ratings_df[ratings_df['tenure'] == 'yes']['beauty'],
                        ratings_df[ratings_df['tenure'] == 'no']['beauty'],
                        center='mean')
```

```
[19]: LeveneResult(statistic=0.48842416527504556, pvalue=0.4849835158609811)
```

Double-click [here](#) for the solution.

1.4.4 Question 4: Using the teachers rating data set, is there an association between visible minorities and tenure?

- Use $\alpha = 0.05$

```
[20]: ## insert code here  
scipy.stats.chi2_contingency(cont_table, correction = True)
```

```
[20]: (2.20678166999886,  
      0.1374050603563787,  
      1,  
      array([[ 42.95896328,  59.04103672],  
             [152.04103672, 208.95896328]]))
```

Double-click [here](#) for a hint.

Double-click [here](#) for the solution.

1.5 Authors

[Aije Egwaikhide](#) is a Data Scientist at IBM who holds a degree in Economics and Statistics from the University of Manitoba and a Post-grad in Business Analytics from St. Lawrence College, Kingston. She is a current employee of IBM where she started as a Junior Data Scientist at the Global Business Services (GBS) in 2018. Her main role was making meaning out of data for their Oil and Gas clients through basic statistics and advanced Machine Learning algorithms. The highlight of her time in GBS was creating a customized end-to-end Machine learning and Statistics solution on optimizing operations in the Oil and Gas wells. She moved to the Cognitive Systems Group as a Senior Data Scientist where she will be providing the team with actionable insights using Data Science techniques and further improve processes through building machine learning solutions. She recently joined the IBM Developer Skills Network group where she brings her real-world experience to the courses she creates.

1.6 Change Log

Date (YYYY-MM-DD)	Version	Changed By	Change Description
2020-08-14	0.1	Aije Egwaikhide	Created the initial version of the lab

Copyright © 2020 IBM Corporation. This notebook and its source code are released under the terms of the [MIT License](#).